

DCR Intelligence Report

Generated: June 20, 2026 at 10:17

Executive Summary

This report summarizes the findings from analyzing 4 agent traces. DecisionAssure discovered 15 unique capabilities, of which 5 are classified as critical. Previously undetected by existing controls: 4 incidents. Current governance coverage: 93.8%.

Metric	Value
Total Traces Analyzed	4
Total Capabilities Discovered	15
Critical Capabilities	5
Previously Missed Incidents	4
Governance Coverage	93.8%
Pending Review	0

Governance Coverage = Verified Capabilities Governed ÷ Total Capabilities Observed x 100. A higher percentage indicates stronger governance maturity.

Key findings indicate that proactive governance updates can significantly reduce risk exposure.

Governance Gap Assessment

Input Assessment

- Traces Analyzed: 4
- Capabilities Discovered: 15
- Previously Undetected Incidents: 4
- Critical Controls Needed: 5

Coverage Improvement

	Before Assessment	After Assessment
Governance Coverage	50.0%	93.8%
Gap		+43.8%

This assessment demonstrates measurable governance improvement.

Executive Impact

4 Previously Missed Incidents Found

Potential Outcomes:

- Evidence Suppression (Severity: high)
- Unauthorized Model Export (Severity: critical)
- Cross-Agent Leakage (Severity: high)
- Shadow Planning (Severity: medium)

Recommended Priority: Immediate review of 5 critical capabilities.

Addressing these gaps reduces the likelihood of credential exposure, model theft, and unauthorized delegation.

Risk Matrix

Severity	Count	Action
Critical	5	Immediate review required
High	6	High priority
Medium	3	Scheduled review
Low	1	Informational

The severity distribution helps prioritize remediation efforts.

Recommended Controls

Critical Severity

- **Unauthorized Model Export:** DENY - export_model_without_approval
- **Approval Bypass:** DENY - bypass_check
- **Model Exfiltration:** DENY - read_model, export_model

High Severity

- **Evidence Suppression:** DENY - delete_logs
- **Cross-Agent Leakage:** DENY - share_data, receive_data
- **Policy Drift:** HUMAN_REVIEW - execute_outdated_policy

Medium Severity

- **Shadow Planning:** MONITOR - plan_secretly
- **Autonomous Tool Chaining:** MONITOR - tool_a, tool_b
- **DCR-2026-0005:** DENY - read_credentials, export_data

Low Severity

- **DCR-2026-0002:** DENY - execute_task, execute_task

Implementing these controls will close governance coverage gaps and reduce risk.

Methodology

1. Discover

Unsupervised clustering identifies recurring action patterns from traces.

2. Verify

Each capability is counterfactually verified - removing any action makes it disappear.

3. Review

Unknown capabilities are routed to human analysts with suggested labels and priority scores.

4. Learn

Approved capabilities are added to the governance ontology, auto-classifying future traces.

5. Replay

The updated ontology is replayed against historical traces to measure coverage improvement.

6. Govern

Governance controls are recommended to close coverage gaps.

This methodology ensures that each discovered capability is verifiable, auditable, and actionable.

Technical note: Discovery uses DBSCAN clustering with TF-IDF vectorisation on trace signatures. Witness hashes are generated using SHA-256. Counterfactual verification removes each required action individually.